

sprint3__autoencoder

December 23, 2022

1 Clustering through feature extraction with an autoencoder

In sprint2 we tried to build a recommendation system for users. A user could load an image and it'll recommend restaurant with similar images. An example is when we input a pasta image that it would recommend restaurants that sell pasta.

In the previous sprint we tried using HOG and SIFT as feature extraction, but that didn't work quite good sadly. Now we'll try to use an autoencoder for feature extraction.

```
[1]: import tensorflow as tf
      from tensorflow import keras
      import os
      import numpy as np
      from fastai.vision.all import PILImage
      import matplotlib.pyplot as plt
```

```
2022-12-22 09:11:53.328672: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: AVX2 AVX512F FMA
```

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
2022-12-22 09:11:53.729778: E tensorflow/stream_executor/cuda/cuda_blas.cc:2981]
Unable to register cuBLAS factory: Attempting to register factory for plugin
cuBLAS when one has already been registered
```

```
2022-12-22 09:11:55.058855: W
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
dynamic library 'libnvinfer.so.7'; dlderror: libnvinfer.so.7: cannot open shared
object file: No such file or directory
```

```
2022-12-22 09:11:55.058996: W
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
dynamic library 'libnvinfer_plugin.so.7'; dlderror: libnvinfer_plugin.so.7:
cannot open shared object file: No such file or directory
```

```
2022-12-22 09:11:55.059010: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Cannot
dlopen some TensorRT libraries. If you would like to use Nvidia GPU with
TensorRT, please make sure the missing libraries mentioned above are installed
properly.
```

1.1 Checking if gpu is available

```
[2]: from tensorflow.python.client import device_lib
print(device_lib.list_local_devices())
```

```
[name: "/device:CPU:0"
device_type: "CPU"
memory_limit: 268435456
locality {
}
incarnation: 3971546288231065025
xla_global_id: -1
, name: "/device:GPU:0"
device_type: "GPU"
memory_limit: 10910760960
locality {
  bus_id: 2
  numa_node: 1
  links {
  }
}
incarnation: 11866661542868266544
physical_device_desc: "device: 0, name: NVIDIA GeForce GTX 1080 Ti, pci bus id:
0000:b0:00.0, compute capability: 6.1"
xla_global_id: 416903419
]
```

2022-12-22 09:12:00.935320: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 AVX512F FMA To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

2022-12-22 09:12:02.218620: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device /device:GPU:0 with 10405 MB memory: -> device: 0, name: NVIDIA GeForce GTX 1080 Ti, pci bus id: 0000:b0:00.0, compute capability: 6.1

1.2 Loading dataset

```
[3]: IMG_HEIGHT = 128
IMG_WIDTH = 128
img_folder = "../tripadvisor_dataset/tripadvisor_images_small"

def create_dataset(img_folder, n=None):
    # n = amount of images
    image_files=os.listdir(os.path.join(img_folder))
    if n==None:
```

```

    n=len(image_files)
    images = np.zeros((n, IMG_HEIGHT, IMG_WIDTH, 3))
    for i,file in enumerate(image_files[:n]):
        # print(f"{i},{file}")
        img=PILImage.create(os.path.join(img_folder,file))
        img_resized=img.resize((IMG_HEIGHT,IMG_WIDTH))
        img_np = np.array(img_resized) #.reshape((IMG_HEIGHT,IMG_WIDTH, 3)) #.
        ↪astype(np.float32)
        images[i]=img_np/255
    return images

images = create_dataset(img_folder)
images_length = len(images)
X_train_full = images
X_train_full.shape

```

[3]: (15182, 128, 128, 3)

Before we continue, we want to filter out ‘bad’ images. ‘bad’ images are images that aren’t food. In the notebook [differentiating-buildings-from-food-cnn.ipynb](#) we’ve made a model to separate food from buildings.

```

[4]: model = keras.models.load_model("../results")
      results = model.predict(X_train_full)

      indices = np.argwhere(results < 0.5)
      indices = indices[..., 0]
      test = X_train_full
      X_train_full = X_train_full[indices]
      X_train_full.shape

```

```

2022-12-22 09:13:17.544655: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 10405 MB memory: -> device:
0, name: NVIDIA GeForce GTX 1080 Ti, pci bus id: 0000:b0:00.0, compute
capability: 6.1
2022-12-22 09:13:36.414850: I tensorflow/stream_executor/cuda/cuda_dnn.cc:384]
Loaded cuDNN version 8500

475/475 [=====] - 14s 24ms/step

```

[4]: (10886, 128, 128, 3)

Now we have filtered out all bad images. We take a look at the bad images to see what the model filtered out, it should normally be all images without food.

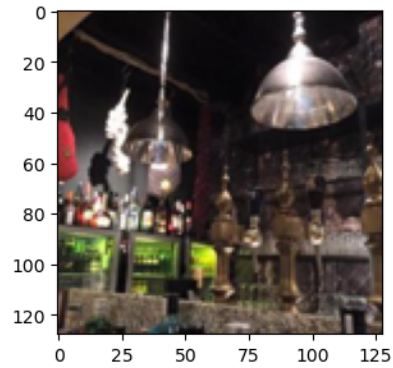
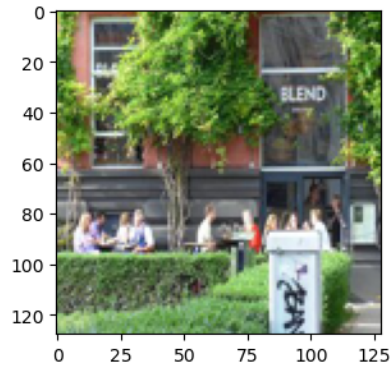
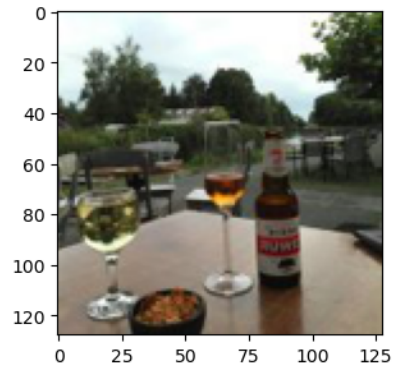
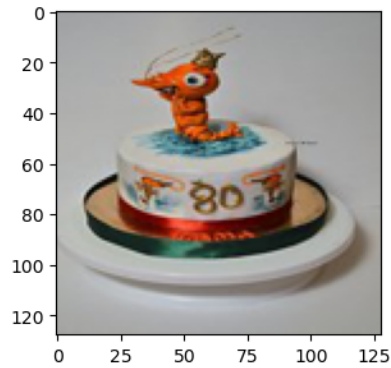
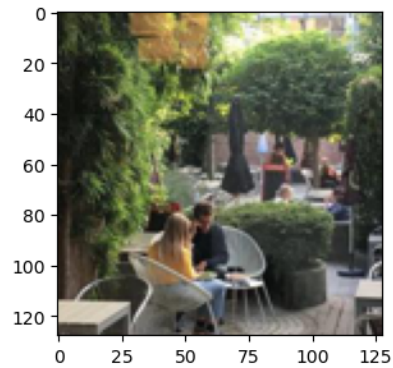
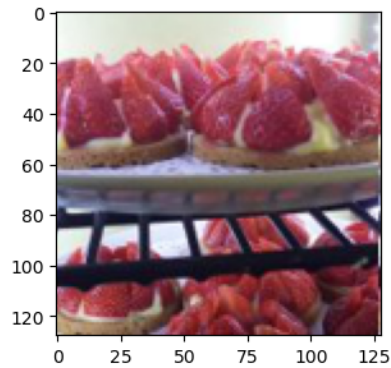



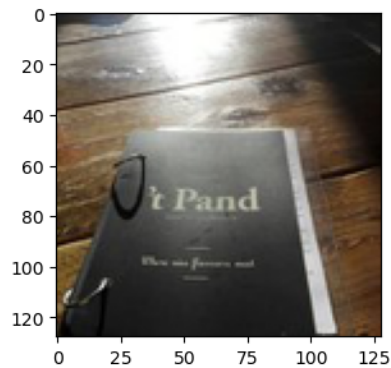
```

[5]: indices_buildings = np.argwhere(results >= 0.5)
      indices_buildings = indices_buildings[..., 0]

```

```
buildings = test[indices_buildings]

fig=plt.figure(figsize=(10,15))
for i in range(0,10):
    plt.subplot(5,2,i+1)
    img = buildings[i]
    # restaurant name get be obtained in file_names[i]
    plt.imshow(img)
fig.tight_layout()
plt.show()
```



That looks good, two images with food are filtered out but the other images are correctly filtered out.

Splitting the dataset in a very small (5 images) test dataset, just to see how the autoencoder will work on new images.

```
[6]: from sklearn.model_selection import train_test_split

positie = len(X_train_full) - 5
test_set = X_train_full[positie:len(X_train_full)]
X_train_full = X_train_full[:positie]

print(test_set.shape)
print(X_train_full.shape)
```

```
(5, 128, 128, 3)
```

```
(10881, 128, 128, 3)
```

1.3 Training the model

The bottleneck of our autoencoder has a shape of (16, 16, 128). This is the shape of the output of the encoder. The ratio between the input and the output of the encoder is $16 \times 16 \times 128 / 128 \times 128 \times 3 = 66.67\%$ (The output features consist of 66.67% of the total amount of pixel values)

```
[7]: def rounded_accuracy(y_true, y_pred):
    return keras.metrics.binary_accuracy(tf.round(y_true), tf.round(y_pred))

conv_encoder = keras.models.Sequential([
    # keras.layers.Reshape([128, 128, 3], input_shape=[128, 128, 3]),
    keras.layers.Conv2D(32, kernel_size=3, padding="SAME", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Conv2D(64, kernel_size=3, padding="SAME", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Conv2D(128, kernel_size=3, padding="SAME", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2)
])
conv_decoder = keras.models.Sequential([
    keras.layers.Conv2DTranspose(64, kernel_size=3, strides=2, padding="SAME",
    ↪activation="selu", input_shape=[16, 16, 128]),
    keras.layers.Conv2DTranspose(32, kernel_size=3, strides=2, padding="SAME",
    ↪activation="selu"),
    keras.layers.Conv2DTranspose(3, kernel_size=3, strides=2, padding="SAME",
    ↪activation="sigmoid"),
])
```

```
conv_ae = keras.models.Sequential([conv_encoder, conv_decoder])

conv_ae.compile(loss="binary_crossentropy", optimizer=keras.optimizers.
    ↪SGD(learning_rate=0.1),
                metrics=[rounded_accuracy])
history = conv_ae.fit(X_train_full, X_train_full, epochs=10, validation_split=0.
    ↪2)

# print(conv_encoder.summary())
# conv_decoder.summary()
```

```
Epoch 1/10
272/272 [=====] - 16s 48ms/step - loss: 0.5767 -
rounded_accuracy: 0.8325 - val_loss: 0.5447 - val_rounded_accuracy: 0.8937
Epoch 2/10
272/272 [=====] - 10s 35ms/step - loss: 0.5409 -
rounded_accuracy: 0.8986 - val_loss: 0.5403 - val_rounded_accuracy: 0.9028
Epoch 3/10
272/272 [=====] - 10s 36ms/step - loss: 0.5366 -
rounded_accuracy: 0.9060 - val_loss: 0.5368 - val_rounded_accuracy: 0.9058
Epoch 4/10
272/272 [=====] - 10s 36ms/step - loss: 0.5344 -
rounded_accuracy: 0.9099 - val_loss: 0.5325 - val_rounded_accuracy: 0.9176
Epoch 5/10
272/272 [=====] - 9s 34ms/step - loss: 0.5323 -
rounded_accuracy: 0.9143 - val_loss: 0.5327 - val_rounded_accuracy: 0.9167
Epoch 6/10
272/272 [=====] - 10s 36ms/step - loss: 0.5311 -
rounded_accuracy: 0.9166 - val_loss: 0.5387 - val_rounded_accuracy: 0.9064
Epoch 7/10
272/272 [=====] - 9s 34ms/step - loss: 0.5300 -
rounded_accuracy: 0.9195 - val_loss: 0.5390 - val_rounded_accuracy: 0.9004
Epoch 8/10
272/272 [=====] - 10s 36ms/step - loss: 0.5289 -
rounded_accuracy: 0.9217 - val_loss: 0.5309 - val_rounded_accuracy: 0.9172
Epoch 9/10
272/272 [=====] - 9s 35ms/step - loss: 0.5285 -
rounded_accuracy: 0.9224 - val_loss: 0.5298 - val_rounded_accuracy: 0.9230
Epoch 10/10
272/272 [=====] - 10s 35ms/step - loss: 0.5277 -
rounded_accuracy: 0.9238 - val_loss: 0.5289 - val_rounded_accuracy: 0.9264
```

Plotting training and loss curve

```
[8]: train_loss_values = history.history['loss']
     val_loss_values = history.history['val_loss']
     best_val_idx = np.argmin(val_loss_values)
```


```


num_epochs = range(len(train_loss_values))

plt.plot(num_epochs, train_loss_values, label='training', color='blue', ls='-')
plt.plot(num_epochs, val_loss_values, label='validatie', color='orange',
         ls='--')
plt.vlines(x=best_val_idx, ymin=0, ymax=1, label='beste validatie',
         color='red', ls=':')
plt.legend()
plt.figure(0)
train_loss_values = history.history['rounded_accuracy']
val_loss_values = history.history['val_rounded_accuracy']
num_epochs = range(len(train_loss_values))

plt.plot(num_epochs, train_loss_values, label='training', color='blue', ls='-')
plt.plot(num_epochs, val_loss_values, label='validatie', color='orange',
         ls='--')
plt.legend()
plt.figure(1)
plt.show()

```





We immediately start with a high accuracy and it keeps increasing. Our loss however almost doesn't decrease.

We check how the test images look when passing through the autoencoder.


```
[9]: def plot_image(image):
    plt.imshow(image, cmap="binary")
    plt.axis("off")

def show_reconstructions(model, images=test_set, n_images=5):
    reconstructions = model.predict(images[:n_images])
    fig = plt.figure(figsize=(n_images * 1.5, 3))
    for image_index in range(n_images):
        plt.subplot(2, n_images, 1 + image_index)
        plot_image(images[image_index])
        plt.subplot(2, n_images, 1 + n_images + image_index)
        plot_image(reconstructions[image_index])

show_reconstructions(conv_ae)

plt.show()
```

1/1 [=====] - 0s 233ms/step



The reconstructed images look pretty good, that means that the feature extraction works. If the images wouldn't look like the original ones, that would mean that the features after encoding aren't representative. In this case the features after encoding are representative because the original image can be almost reconstructed.

```
[10]: X_train_full[:1].shape
```

```
[10]: (1, 128, 128, 3)
```

1.4 Save model

We only save the encoder, we don't need the decoder for feature extraction. The decoder was only needed when training.

```
[11]: conv_encoder.save('./model_last_version/')
```

```
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.
```

```
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 3 of 3). These functions will not be directly callable after loading.
```

```
INFO:tensorflow:Assets written to: ./model_last_version/assets
```

```
INFO:tensorflow:Assets written to: ./model_last_version/assets
```

1.5 Getting features of all the images

These features will be used in another notebook.

```
[12]: features = conv_encoder.predict(X_train_full)
```

```
341/341 [=====] - 3s 9ms/step
```

```
[13]: import pickle
      with open('features_encoder_pictures_with_no_buildings_last_version.pkl', 'wb') as f:
          pickle.dump(features, f)
```

In notebook [sprint3_autoencoder_features.ipynb](#) we will use these features saved to the pickle file to cluster

1.6 References

<https://github.com/ageron/handson-ml2>